

Machine Learning

Homework 1

RAPPAPORT Gabrielle

September 14, 2019

1 Question 1

We are given the following constraints :

$$\forall i \in \{1, \dots, n\}, a_i^T x \geq b_i$$

We assume that :

$$\exists M : a_i^T x \geq M \text{ and } \forall i \in \{1, \dots, n\}, b_i \geq M$$

We want to model the requirement that at least k of the constraints are satisfied.

Let's introduce the binary variable :

$$\forall i \in \{1, \dots, n\}, y_i \in \{0, 1\} \text{ such as } \sum_{i=1}^n y_i \leq k$$

We can reformulate our constraints as :

$$\forall i \in \{1, \dots, n\}, a_i^T x \geq b_i(1 - y_i) + M.y_i$$

This means that :

$$y_i = 0 \Rightarrow a_i^T x \geq b_i, \text{ the } i^e \text{ constraint is satisfied}$$

$$y_i = 1 \Rightarrow a_i^T x \geq M, \text{ the } i^e \text{ constraint is not satisfied}$$

Finally, we obtain the following model :

$$a_i^T x \geq b_i(1 - y_i) + M.y_i, \forall i \in \{1, \dots, n\}$$

$$\sum_{i=1}^n y_i \leq k$$

$$y_i \in \{0, 1\}$$

2 Question 2

2.1 Question a

We want to solve the classical least square problem :

$$\min \sum_{i=1}^n (y_i - \beta_0 + \beta^T x_i)^2$$

We name β_0^* and β^* the optimal solutions.

To solve this problem, we will use the matrix representation. We name :

$$\bar{\beta} = \begin{pmatrix} \beta_0^* \\ \beta^* \end{pmatrix} \text{ and } \bar{X} = \begin{pmatrix} 1 & x_1 & \cdots & x_d \end{pmatrix}$$

Which give us the equivalent problem :

$$\begin{aligned} \sum_{i=1}^n (y_i - \beta_0 - \beta^T x_i)^2 &= \|y - \bar{X} \cdot \bar{\beta}\|_2^2 \\ &= (y - \bar{X} \bar{\beta})^T (y - \bar{X} \bar{\beta}) \end{aligned}$$

To find the minimum we find B which cancels the derivative.

$$\begin{aligned} &\frac{\partial}{\partial \bar{\beta}} \{ (y - \bar{X} \bar{\beta})^T (y - \bar{X} \bar{\beta}) \} \\ &= \frac{\partial}{\partial \bar{\beta}} \{ (y^T - (\bar{X} \bar{\beta})^T) (y - \bar{X} \bar{\beta}) \} \\ &= \frac{\partial}{\partial \bar{\beta}} \{ y^T y - y^T \bar{X} \bar{\beta} - (\bar{X} \bar{\beta})^T y + \bar{\beta}^T \bar{X}^T \bar{X} \bar{\beta} \} \\ &= -2 \bar{X}^T y + 2 \bar{X}^T \bar{X} \bar{\beta} \\ &= -2 \bar{X}^T (y - \bar{X} \bar{\beta}) \end{aligned}$$

Then :

$$-2 \bar{X}^T (y - \bar{X} \bar{\beta}) = 0 \Leftrightarrow \bar{\beta} = \left(\bar{X}^T \bar{X} \right)^{-1} \bar{X}^T y$$

2.2 Question b

We want to solve the following problem effectively :

$$\min \sum_{i=1}^m |y_i - \beta_0 - \beta^T x_i|$$

Since we are adding positive elements, we know that :

$$\min \sum_{i=1}^m |y_i - \beta_0 - \beta^T x_i| = \sum_{i=1}^m \min |y_i - \beta_0 - \beta^T x_i|$$

We know that :

$$\min |y_i - \beta_0 - \beta^\top x_i| \Leftrightarrow$$

$$\begin{array}{ll} \min & z_i \\ \text{s.t.} & \left| \begin{array}{l} z_i \geq y_i - \beta_0 - \beta^\top x_i \\ z_i \geq -y_i + \beta_0 + \beta^\top x_i \end{array} \right. \end{array}$$

This finally gives us the linear program, which is the most efficient way to solve the problem :

$$\begin{array}{ll} \min & \sum_{i=1}^n z_i \\ \text{s.t.} & \left| \begin{array}{l} z_i \geq y_i - \beta_0 - \beta^\top x_i, \forall i \\ z_i \geq -y_i + \beta_0 + \beta^\top x_i, \forall i \end{array} \right. \end{array}$$

2.3 Question c

In this question, we want to model that at most k of the d coefficient β_1, \dots, β_d have non zero values.

To do so, we will use the *big-M method*. We introduce, a constant M big enough and new **binary** variables a_i such as :

$$\beta_i \neq 0 \Rightarrow a_i = 1$$

We obtain the following linear problem :

$$\begin{array}{ll} \min & \sum_{i=1}^n z_i \\ \text{s.t.} & \left| \begin{array}{l} z_i \geq y_i - \beta_0 - \beta^\top x_i, \forall i \\ z_i \geq -y_i + \beta_0 + \beta^\top x_i, \forall i \\ \beta_i \leq M.a_i \\ \beta_i \geq -M.a_i \\ \sum_{i=1}^d a_i \leq k \\ a_i \in \{0,1\}, \forall i \end{array} \right. \end{array}$$

2.4 Question d

We want to formulate the *Robust Linear Regression Problem* as an integer programming problem:

$$\begin{aligned} \min \text{Median}(|y_1 - \beta_0 - \beta^T x_1|, \dots, |y_n - \beta_0 - \beta^T x_n|) \\ \text{where } n = 2k + 1 \end{aligned}$$

Since the number of elements $|y_i - \beta_0 - \beta^T x_i|$ is odd, the median is $|y_{k+1} - \beta_0 - \beta^T x_{k+1}|$, therefore we want to minimize $c = |y_{k+1} - \beta_0 - \beta^T x_{k+1}|$.

It means that we want to minimize c such as c is bigger than the first $k + 1$ elements.

This gives us the following linear program :

$$\begin{aligned} \min c \\ \text{s.t.} \quad & \left| \begin{array}{ll} c & \geq h_i - M(1 - z_i), \forall i \\ h_i & \geq y_i - \beta_0 - \beta^T x_i, \forall i \\ h_i & \geq -y_i + \beta_0 + \beta^T x_i, \forall i \\ \sum_{i=1}^n z_i & = k \\ z_i & \in \{0,1\}, \forall i \end{array} \right. \end{aligned} \quad \text{ATTENTION : } k+1 \rightarrow \text{Les } k+1 \text{ contraintes sont vérifiées.}$$

3 Question 3

September 14, 2019

1 Exercise 3

```
[2]: #import Pkg
      #Pkg.add("JuMP")
      #Pkg.add("Gurobi")

      using JuMP, Gurobi, DataFrames, CSV, Random, LinearAlgebra

[3]: #using Pkg
      #Pkg.add("PyPlot")
      using PyPlot
```

2 Question a

- Model for the l_0 -regularization

```
[4]: function regularized_regression_l0(y, X, ρ)

      M = 10000

      n,m = size(X)
      model = Model(solver = GurobiSolver(TimeLimit=45))

      @variable(model, β[1:m])
      @variable(model, z[1:m], Bin)

      @objective(model, Min, sum((y - X*β).^2) + ρ*sum(z))

      for i in 1:m
          @constraint(model, β[i] <= M * z[i])
          @constraint(model, (-β[i]) <= M * z[i])
      end

      sol = solve(model)

      return (getobjectivevalue(model), getvalue(β))

end
```

[4]: regularized_regression_l0 (generic function with 1 method)

- Model for the l1-regularization

```
[6]: function regularized_regression_l1(y, X,  $\rho$ )

    M = 10000

    n,m = size(X)
    model = Model(solver = GurobiSolver(TimeLimit=45))

    @variable(model,  $\beta$ [1:m])
    @variable(model, z[1:m])

    @objective(model, Min, sum((y - X* $\beta$ ).^2) +  $\rho$ *sum(z))

    for i in 1:m
        @constraint(model, (z[i]) >=  $\beta$ [i])
        @constraint(model, (z[i]) >= -  $\beta$ [i])
    end

    sol = solve(model)

    return (getobjectivevalue(model), getvalue( $\beta$ ))

end
```

[6]: regularized_regression_l1 (generic function with 1 method)

- Model for the l2-regularization

```
[7]: function regularized_regression_l2(y, X,  $\rho$ )

    M = 10000

    n,m = size(X)
    model = Model(solver = GurobiSolver(TimeLimit=45))

    @variable(model,  $\beta$ [1:m])

    @objective(model, Min, sum((y - X* $\beta$ ).^2) +  $\rho$ *sum( $\beta$ .^2))

    sol = solve(model)

    return (getobjectivevalue(model), getvalue( $\beta$ ))

end
```

[7]: regularized_regression_l2 (generic function with 1 method)

2.1 Question b

```
[ ]: sparseX2 = CSV.read("sparseX2.csv")
     sparseY2 = CSV.read("sparseY2.csv")
```

Creation of the validation, training and testing sets.

```
[9]: function split_data(X, y, val, test)
      n = size(X, 1)
      index = shuffle([i for i in 1:n])

      size_validation = floor(Int, val*n)
      size_test = floor(Int, (val+test)*n)

      ind_validation = index[1:size_validation]
      ind_test = index[size_validation+1:size_test]
      ind_train = index[size_test+1:n]

      X_validation = X[ind_validation, :]
      X_test = X[ind_test, :]
      X_train = X[ind_train, :]

      y_validation = y[ind_validation, :]
      y_test = y[ind_test, :]
      y_train = y[ind_train, :]

      return (
        convert(Matrix, X_validation),
        convert(Matrix, X_test),
        convert(Matrix, X_train),
        convert(Matrix, y_validation),
        convert(Matrix, y_test),
        convert(Matrix, y_train)
      )

end
```

[9]: split_data (generic function with 1 method)

```
[ ]: (
      X_validation,
      X_test,
      X_train,
      y_validation,
      y_test,
      y_train
    ) = split_data(sparseX2, sparseY2, 0.25, 0.25)
```

We choose ρ as the value that gives the best mean squared prediction error on the validation set.

```
[12]: function mean_square_error(y_actual, y_predicted)
      return sum((y_actual - y_predicted).^2)
end
```

[12]: mean_square_error (generic function with 1 method)

```
[13]: function compute_error_l0( $\rho$ , X, y)
      objective,  $\beta$  = regularized_regression_l0(y_train, X_train,  $\rho$ )
      error = mean_square_error(y, X *  $\beta$ )
      return error
end
```

[13]: compute_error_l0 (generic function with 1 method)

```
[14]: function compute_error_l1( $\rho$ , X, y)
      objective,  $\beta$  = regularized_regression_l1(y_train, X_train,  $\rho$ )
      error = mean_square_error(y, X *  $\beta$ )
      return error
end
```

[14]: compute_error_l1 (generic function with 1 method)


```
[15]: function compute_error_l2( $\rho$ , X, y)
      objective,  $\beta$  = regularized_regression_l2(y_train, X_train,  $\rho$ )
      error = mean_square_error(y, X *  $\beta$ )
      return error
end
```

[15]: compute_error_l2 (generic function with 1 method)

```
[ ]:  $\rho_{10}$  = vcat([0.001, 0.01, 0.1], [i for i=1:30])
      errors_l0 = zeros(length( $\rho_{10}$ ))

      for i in 1:length(errors_l0)
          errors_l0[i] = compute_error_l0( $\rho_{10}$ [i], X_validation, y_validation)
      end
```

```
[17]: plot( $\rho_{10}$ , errors_l0)
```

```
[17]: 1-element Array{PyCall.PyObject,1}:
      PyObject <matplotlib.lines.Line2D object at 0x14212a748>
```


```
[41]:  $\rho_{\text{optimal\_10}} = \rho_{\text{10}}[\text{argmin}(\text{errors\_10})]$ 
```

```
[41]: 30.0
```

```
[ ]:  $\rho_{\text{11}} = [i/2 \text{ for } i=1:60]$ 
      errors_11 = zeros(length( $\rho_{\text{11}}$ ))

      for i in 1:length(errors_11)
          errors_11[i] = compute_error_11( $\rho_{\text{11}}[i]$ , X_validation, y_validation)
      end
```

```
[20]: plot( $\rho_{\text{11}}$ , errors_11)
```



```
[20]: 1-element Array{PyCall.PyObject,1}:
      PyObject <matplotlib.lines.Line2D object at 0x1428417b8>
```


```
[42]:  $\rho_{\text{optimal\_l1}}$  =  $\rho_{\text{l1}}[\text{argmin}(\text{errors\_l1})]$ 
```

```
[42]: 10.0
```

```
[ ]:  $\rho_{\text{l2}}$  = vcat([0.001, 0.01, 0.1], [i for i=1:20])
      errors_l2 = zeros(length( $\rho_{\text{l2}}$ ))

      for i in 1:length(errors_l2)
          errors_l2[i] = compute_error_l2( $\rho_{\text{l2}}[i]$ , X_validation, y_validation)
      end
```

```
[28]: plot( $\rho_{\text{l2}}$ , errors_l2)
```



```
[28]: 1-element Array{PyCall1.PyObject,1}:
      PyObject <matplotlib.lines.Line2D object at 0x141a41f60>
```

```
[43]: ρ_optimal_l2 = ρ_l2[argmin(errors_l2)]
```

```
[43]: 4.0
```

Compute $\|y - X\beta\|_2$ on the testing set using the β 's from l0-regularized, l1-regularized linear regression, l2-regularized linear regression, and standard linear regression ($\rho = 0$)

```
[50]: l0_error = compute_error_l0(ρ_optimal_l0, X_test, y_test)
```

Academic license - for non-commercial use only

Optimize a model with 200 rows, 200 columns and 400 nonzeros

Model has 5050 quadratic objective terms

Variable types: 100 continuous, 100 integer (100 binary)

Coefficient statistics:

Matrix range [1e+00, 1e+04]

Objective range [3e+00, 3e+02]

QObjective range [2e-02, 3e+02]

Bounds range [1e+00, 1e+00]

RHS range [0e+00, 0e+00]

Found heuristic solution: objective 0.0000000

Presolve time: 0.00s

Presolved: 200 rows, 200 columns, 400 nonzeros

Presolved model has 5050 quadratic objective terms
 Variable types: 100 continuous, 100 integer (100 binary)

Root relaxation: objective -9.524350e+02, 981 iterations, 0.03 seconds

Nodes		Current Node			Objective Bounds			Work	
Expl	Unexpl	Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
	0	0	-952.43503	0	78	0.00000 -952.43503	-	-	0s
H	0	0				-309.6583805 -952.43503	208%	-	0s
H	0	0				-413.9725003 -952.43503	130%	-	0s
	0	0	-952.43503	0	78	-413.97250 -952.43503	130%	-	0s
H	0	0				-631.5297179 -952.43503	50.8%	-	0s
H	0	0				-631.8148568 -952.43503	50.7%	-	0s
	0	2	-952.43503	0	78	-631.81486 -952.43503	50.7%	-	0s
H	744	552				-632.3423572 -884.06346	39.8%	27.2	1s
H	2792	1173				-632.3758008 -847.67096	34.0%	28.1	3s
H	2810	1186				-632.4516075 -846.68895	33.9%	28.1	3s
H	2816	1190				-633.1807952 -846.68895	33.7%	28.1	3s
	5002	2258	-725.62254	28	75	-633.18080 -828.40078	30.8%	28.1	5s
H	5971	2722				-633.3912979 -822.42464	29.8%	28.0	5s
	12034	5283	-670.01060	32	72	-633.39130 -799.96555	26.3%	27.8	10s
H1	14090	6176				-633.5924643 -794.95228	25.5%	28.0	12s
	17025	7198	-706.92109	28	77	-633.59246 -788.03072	24.4%	27.9	15s
	24809	9817	-741.81739	26	79	-633.59246 -774.31503	22.2%	27.5	20s
H2	27145	10533				-633.6506843 -771.17807	21.7%	27.4	21s
	31627	11967	-714.26739	27	74	-633.65068 -766.26198	20.9%	27.3	25s
	39822	14334	-682.52802	28	69	-633.65068 -758.31175	19.7%	27.3	30s
H4	40251	14434				-633.9829945 -757.89283	19.5%	27.3	30s
H4	40311	14393				-634.6981061 -757.78877	19.4%	27.2	30s
	47551	16181	-706.87827	30	75	-634.69811 -751.93922	18.5%	27.1	35s
	54518	17681	-746.29602	27	78	-634.69811 -746.33075	17.6%	27.0	40s
	61880	19284	-650.50638	34	68	-634.69811 -741.44961	16.8%	26.9	45s

Explored 62113 nodes (1674274 simplex iterations) in 45.00 seconds

Thread count was 4 (of 4 available processors)

Solution count 10: -634.698 -633.983 -633.651 ... -631.815

Time limit reached

Best objective -6.346981061123e+02, best bound -7.413865505370e+02, gap 16.8093%

[U+250C] Warning: Not solved to optimality, status: UserLimit

[U+2514] @ JuMP /Users/gabriellerappaport/.julia/packages/JuMP/I7whV/src/solvers.

↪jl:212

[50]: 16.229156793039543

```
[53]: l1_error = compute_error_l1( $\rho$ _optimal_l1, X_test, y_test)
```

```
Academic license - for non-commercial use only
Optimize a model with 200 rows, 200 columns and 400 nonzeros
Model has 5050 quadratic objective terms
Coefficient statistics:
  Matrix range      [1e+00, 1e+00]
  Objective range   [3e+00, 3e+02]
  QObjective range  [2e-02, 3e+02]
  Bounds range      [0e+00, 0e+00]
  RHS range         [0e+00, 0e+00]
Presolve removed 100 rows and 0 columns
Presolve time: 0.00s
Presolved: 100 rows, 200 columns, 200 nonzeros
Presolved model has 5050 quadratic objective terms
Ordering time: 0.00s
```

```
Barrier statistics:
Free vars   : 199
AA' NZ      : 9.900e+03
Factor NZ   : 1.219e+04
Factor Ops  : 1.092e+06 (less than 1 second per iteration)
Threads     : 1
```

Iter	Objective		Residual		Compl	Time
	Primal	Dual	Primal	Dual		
0	2.14000000e+06	0.00000000e+00	0.00e+00	2.12e+03	1.01e+06	0s
1	2.07540565e+06	-9.51734680e+02	2.95e-07	2.12e-03	1.04e+04	0s
2	2.20304958e+03	-9.51573018e+02	1.57e-10	1.10e-06	1.58e+01	0s
3	-7.31306591e+02	-8.85413674e+02	2.99e-09	1.22e-07	7.71e-01	0s
4	-8.25442948e+02	-8.40833167e+02	4.43e-10	2.05e-08	7.70e-02	0s
5	-8.35758966e+02	-8.37264584e+02	1.84e-10	2.22e-09	7.53e-03	0s
6	-8.36690413e+02	-8.36776719e+02	2.40e-10	2.70e-09	4.32e-04	0s
7	-8.36734721e+02	-8.36739711e+02	2.97e-10	4.46e-09	2.49e-05	0s
8	-8.36736224e+02	-8.36736451e+02	1.49e-10	1.39e-09	1.13e-06	0s
9	-8.36736252e+02	-8.36736254e+02	1.53e-11	2.18e-10	5.73e-09	0s

```
Barrier solved model in 9 iterations and 0.02 seconds
Optimal objective -8.36736252e+02
```

```
[53]: 20.376884563437116
```

```
[47]: l2_error = compute_error_l2( $\rho$ _optimal_l2, X_test, y_test)
```

```
Academic license - for non-commercial use only
Optimize a model with 0 rows, 100 columns and 0 nonzeros
```

Model has 5050 quadratic objective terms
Coefficient statistics:
Matrix range [0e+00, 0e+00]
Objective range [3e+00, 3e+02]
QObjective range [2e-02, 3e+02]
Bounds range [0e+00, 0e+00]
RHS range [0e+00, 0e+00]
Presolve time: 0.00s
Presolved: 0 rows, 100 columns, 0 nonzeros
Presolved model has 5050 quadratic objective terms
Ordering time: 0.00s

Barrier statistics:
Free vars : 199
AA' NZ : 4.851e+03
Factor NZ : 4.950e+03
Factor Ops : 3.284e+05 (less than 1 second per iteration)
Threads : 1

Iter	Objective		Residual		Compl	Time
	Primal	Dual	Primal	Dual		
0	0.00000000e+00	0.00000000e+00	0.00e+00	2.40e+02	0.00e+00	0s
1	-2.79969317e+02	-2.54327433e+01	3.26e-08	2.00e+02	0.00e+00	0s
2	-5.91843984e+02	-1.50106285e+02	4.64e-08	1.43e+02	0.00e+00	0s
3	-8.22895664e+02	-4.23828169e+02	6.84e-08	7.68e+01	0.00e+00	0s
4	-9.16833998e+02	-9.16833396e+02	9.32e-08	7.70e-05	0.00e+00	0s
5	-9.16833985e+02	-9.16833985e+02	6.72e-14	7.72e-11	0.00e+00	0s

Barrier solved model in 5 iterations and 0.01 seconds
Optimal objective -9.16833985e+02

[47]: 180.90945992777353

```
[64]: println(" 10 error: ", 10_error)
println(" 11 error: ", 11_error)
println(" 12 error: ", 12_error)
```

10 error: 16.229156793039543	rho = 3, error 42 ordinary least square : rho = 0 baseline : avg(concat (train validation)
11 error: 20.376884563437116	
12 error: 180.90945992777353	